Robust Learning Under Label Noise for Real-World Classification

Evaluation & Reproducibility Report | Experiment ID: a1d313f3-59f3-4015-9321-d0bdaa96ce81 | Generated: 2026-10-25 14:31:20

1. Dataset & Experimental Setup

Target Variable:	species	**Classes:**	3 (Setosa, Versicolor, Virginica)
Total Rows:	150	**Preprocessed Features:**	4
Train / Test Split:	105 (70%) / 22 (15%)	**Noise Levels Evaluated:**	0%, 10%, 20%, 30%
Random Seed:	[42, 123, 456]	**Noise Strategy:**	Symmetric

2. Model Evaluation Across Noise Levels

Evaluated strictly on the clean holdout test set to measure genuine generalization ability under synthetic training noise.

Model	Type	Noise	Accuracy	Precision	Recall	F1-Score	Time (s)
Baseline Logistic Regression	Baseline	0%	0.8182	0.8303	0.8182	**0.8144**	0.01
Baseline Random Forest	Baseline	0%	0.9091	0.9091	0.9091	**0.9091**	0.46
Robust Sample-Weighted (Random Forest)	Robust	0%	0.9091	0.9091	0.9091	**0.9091**	0.33
Robust Trimmed (Random Forest)	Robust	0%	0.9091	0.9091	0.9091	**0.9091**	0.24
Baseline Logistic Regression	Baseline	10%	0.8636	0.8674	0.8636	**0.8629**	0.02
Baseline Random Forest	Baseline	10%	0.9091	0.9091	0.9091	**0.9091**	0.34
Robust Sample-Weighted (Random Forest)	Robust	10%	0.9091	0.9091	0.9091	**0.9091**	0.45
Robust Trimmed (Random Forest)	Robust	10%	0.9091	0.9091	0.9091	**0.9091**	0.38
Baseline Logistic Regression	Baseline	20%	0.8182	0.8303	0.8182	**0.8144**	0.01
Baseline Random Forest	Baseline	20%	0.8182	0.8295	0.8182	**0.8212**	0.41
Robust Sample-Weighted (Random Forest)	Robust	20%	0.8182	0.8295	0.8182	**0.8212**	0.36
Robust Trimmed (Random Forest)	Robust	20%	0.8636	0.8674	0.8636	**0.8629**	0.34
Baseline Logistic Regression	Baseline	30%	0.8636	0.9045	0.8636	**0.8571**	0.02
Baseline Random Forest	Baseline	30%	0.7727	0.7822	0.7727	**0.7751**	0.37
Robust Sample-Weighted (Random Forest)	Robust	30%	0.7727	0.7822	0.7727	**0.7751**	0.34
Robust Trimmed (Random Forest)	Robust	30%	0.8182	0.8295	0.8182	**0.8212**	0.34

3. Confident Learning Label Issue Detection

Cross-validation out-of-fold probability estimation identified **11** suspicious training samples (10.48% of training split) exceeding the issue threshold of 0.50. These samples exhibited substantial margin disagreement between the assigned label and the learned distribution.

Row ID	Observed Label	Predicted Class	Confidence	Observed Prob	Issue Score	Status
0	Virginica	Setosa	0.64	0.05	**0.77**	Potential Issue
39	Virginica	Setosa	0.59	0.06	**0.73**	Potential Issue
129	Setosa	Virginica	0.55	0.06	**0.72**	Potential Issue
23	Virginica	Setosa	0.56	0.10	**0.68**	Potential Issue
32	Versicolor	Setosa	0.75	0.21	**0.67**	Potential Issue
127	Setosa	Virginica	0.51	0.11	**0.65**	Potential Issue
98	Virginica	Versicolor	0.56	0.18	**0.60**	Potential Issue
71	Setosa	Versicolor	0.51	0.18	**0.58**	Potential Issue

4. Empirical Findings & Discussion

• **Baseline Performance:** Under pristine labels (0% noise), Baseline Random Forest achieved the highest F1-score of 0.9091 (Accuracy: 0.9091).

• **Noise Robustness:** At 30% label noise, the conventional model (Baseline Logistic Regression) retained higher F1-score (0.8571) than the robust model (Robust Trimmed (Random Forest), 0.8212). This suggests either the dataset distribution possesses naturally resilient decision margins or the filtering threshold was conservative.

5. Methodological Notes & Reproducibility

Reproducibility: Results are strictly deterministic given the configured random seed and split ratios. Label noise was applied solely to the training set; test labels remained unaltered to ensure valid evaluation.
Technical Note: Confident learning scores indicate statistical discrepancy between observed labels and feature density; they represent potential label noise, not mathematical proof of label error in uncontrolled environments.